

# CPA project

Edwin Ansari  
Maria Sidko

## 1 Introduction

Real world implementation of the mathematical notion of graph theory has countless applications in modern computer science, which have become crucial in our daily life, from efficient handling of data to pertinent responses of search engines. Throughout this project, we have attempted to understand, implement and optimize some of those unavoidable algorithms, following the guidelines provided by our professor for the subject.

In order to reach that goal, we have decided to implement our code using the Java language, given its reasonable calculation efficiency, in addition to its object oriented paradigm and provided APIs, reaching an optimal middle ground between speed and coding comfort.

## 2 TME 1 - Handling a large graph

**Exercise 1 - Preparation** The provided files are comprised of the edges of the graph, so we can easily count the edges by looping through it. The nodes on the other hand should only be counted once, so we used an HashSet for that purpose. The class `Counter` implements this process. Following are the results of the counting on the examples :

|                          |                    |           |
|--------------------------|--------------------|-----------|
| Amazon benchmark :       | number of vertices | 334863    |
|                          | number of edges    | 925872    |
| Live journal benchmark : | number of vertices | 3997962   |
|                          | number of edges    | 34681189  |
| Orkut benchmark :        | number of vertices | 3072441   |
|                          | number of edges    | 117185083 |

N. B. The last provided benchmark, Friendster, which has 65608366 Nodes and 1806067135 Edges according to the website it was taken from. That gives us  $|E| \approx 10^9$  and  $|G| \approx 10^9 + 10^7$ . Just reading this file and storing it in a simple list structure would take  $10^9 \times \text{sizeof(int)} = 10^9 \times 4\text{Bytes} \approx 4\text{GB}$  of RAM and then the whole program that counts it would take on top of it around 5GB of RAM (that is nearly impossible with our basic computers having 8GB of RAM in total). The same reasoning could be applied in terms of time complexity. We were therefore unfortunately unable to use that benchmark throughout the project.

**Exercise 2 - Three graph data structures** Given that the graphs in the provided files are essentially already stored as a list of edges, the parsing, done by the `GraphLoader Class`, creates automatically a list of edges representation (`GraphList Class`), from which are derived the two other representations, `GraphArray Class` and `GraphMatrix Class`. The number of vertices of the two latter representations is `MAX_NODE_ID + 1`. In an attempt to improve speed, we decided to replace java Scanner with a buffered `FastReader.java`. Our results were the following :

|                |  |                   |
|----------------|--|-------------------|
| Amazon :       | List of edges                              | 1934ms            |
|                | List of edges using FastReader             | 939ms             |
|                | Adjacency array                            | 89ms              |
|                | Adjacency matrix                           | OutOfMemoryError  |
| Live-journal : | List of edges                              | 91363ms           |
|                | List of edges (using FastReader)           | 31400ms           |
|                | Adjacency array                            | OutOfMemoryErrorr |
|                | List of edges (allocating 4500MB of RAM)   | 32031ms           |
|                | Adjacency array (allocating 4500MB of RAM) | 17712ms           |
| Orkut :        | List of edges (using FastReader)           | OutOfMemoryError  |
|                | Adjacency array                            | OutOfMemoryErrorr |
|                | List of edges (allocating 5800MB of RAM)   | 441201ms          |
|                | Adjacency array (allocating 5800MB of RAM) | OutOfMemoryErrors |

### Space Complexity :

List of Edges :  $O(|E|)$

Adjacency Array :  $O(|V| + |E|)$

Adjacency Matrix :  $O(|V|^2)$

We have the same time complexity for structures reading.

It is noticeable right away that, while the list of edges representation is scaled up reasonably easily, and so is the adjacency array representation, the matrix representation quickly gets out of range, and despite additional memory allocation, reaches the hardware limits of our machines.

**Exercise 3 - BFS** Diameter results using `GraphDiameter` class :

|              |   |         |
|--------------|---|---------|
| Amazon       | Lower-Bound                                   | 47      |
|              | Lower-Bound time                              | 2330ms  |
|              | Upper-Bound                                   | 57      |
|              | Upper-Bound time (after Lower Bound executed) | 296ms   |
| Live-Journal | Lower-Bound                                   | 21      |
|              | Lower-Bound time                              | 74739ms |
|              | Upper-Bound                                   | 24      |
|              | Upper-Bound time (after Lower Bound executed) | 10685ms |

**Exercise 4 - Triangles** The intuitive way to calculate the triangles would be for each edge of the graph to compute the intersection of each two nodes. However, due to efficiency issues, some adjustments can be made to that algorithm. In order to avoid registering the same triangle multiple times, as well as to optimize processing time, we can visit the vertices in a non increasing degree order, and only considering neighbor vertices of lower degree. This is implemented by the `TriangleFinder` class

|        |               |        |
|--------|---------------|--------|
| amazon | triangles     | 667129 |
|        | counting time | 2840ms |

## 3 TME2

**Exercise 1 - PageRank (Directed Graph)** Implementing the Pagerank algorithm in the `PageRank` class we could achieve accurate results after about 5 or 6 iterations, using the Wikipedia database. Here are the top 5 listed pages :

```
max 1 : United States
max 2 : United Kingdom
max 3 : 2007
max 4 : Germany
max 5 : 2006
```

and the 5 bottom :

```
min 1 : Schismogenesis
min 2 : Thomas W. Thompson
min 3 : Thomas W. Thompson
min 4 : List of examples in general topology
min 5 : James Dempsey
```

**Exercise 2 - Correlations** In order to plot we used linear scales.




FIGURE 1 – Scatter plot using the Wkikipedia dataset

Intuitively, the more references(In Degree) a node has, the higher its rank is. We can notice that on the scatter plot, as it is somehow linearly correlated to the Pagerank results, which means that it is linked to the ranking, but not identical, given that other factors enter in the Pagerank calculation.




FIGURE 2 – Scatter plot using the Wkikipedia dataset




FIGURE 3 – Scatter plot using the Wkikipedia dataset




FIGURE 4 – Scatter plot using the Wkikipedia dataset




FIGURE 5 – Scatter plot using the Wkikipedia dataset




FIGURE 6 – Scatter plot using the Wikipedia dataset

We can notice that the  $\alpha$  value of the PageRank algorithm has an influence on its results, the closer they are the more linear the plot is, meaning that the PageRank results of the two calls are similar. When the values become more distant, the rankings change, leaving the scatter plot more sparse.

### Exercise 3 - Personalized Page Rank

**Magnus Carlsen** : To find very close results to node(Magnus Carlsen) we first found its Page Id : 442682 and then run the *rootedPageRank* using an initial  $P_0$  vector such that :

$$P_0(\text{magnus\_id}) = 1$$

$$P_0(i) = 0 \quad \forall i \neq \text{magnus\_id}$$

By computing personalized PageRank using *Magnus Carlsen* as the root node, we obtained the following top and bottom pages :

```
max 1 : Magnus Carlsen
max 2 : Chess
max 3 : Grandmaster (chess)
max 4 : Federation Internationale des echecs
max 5 : Elo rating system

min 1 : Transmission and infection of H5N1
min 2 : Ital
min 3 : (Sittin' on) the Dock of the Bay
min 4 : Ministry of New and Renewable Energy (India)
min 5 : Superclub
```

#### Chess Boxing :

In order to have a relevant restart vector using a set of subjects (Chess and Boxing in this case) we offer two methods.

1. In this first method we simply set the restarting vector  $P_0$  such that :

$$P_0(\text{id}) = \frac{1}{N} \text{ if } \text{id} \in S$$

$$P_0(\text{id}) = 0 \quad \forall \text{id} \notin S$$

Where  $S$  is the set of pages that are linked to any of the related categories and  $N$  is their quantity.

The results using this method are :

max 1 : Chess  
 max 2 : Grandmaster (chess)  
 max 3 : Federation Internationale des echecs  
 max 4 : Chess Olympiad  
 max 5 : United States  
 max 6 : Boxing  
 max 7 : International Master  
 max 8 : Elo rating system  
 max 9 : World Chess Championship  
 max 10 : Chess opening

min 1 : Villeray–Saint–Michel–Parc–Extension  
 min 2 : Jessica Cutler  
 min 3 : Electrical connection  
 min 4 : Bass player  
 min 5 : Caribou–Targhee National Forest  
 min 6 : Above the Rim  
 min 7 : Charlesbourg , Quebec  
 min 8 : Gaius Calpurnius Piso  
 min 9 : Lisa Opie  
 min 10 : Persistence

2. We fetch, using all the provided files, the number of categories related to the requested subjects each page is linked to, alongside with searching whether the actual names of the pages contain those subjects. We use this information to weight the restarting vector  $P_0$  such that :  $P_0(i) = \frac{r+10}{t}$  if  $ct = 0$   
 $P_0(i) = \frac{r+ct*k}{t}$  if  $ct > 0$   
 $P_0(i) = 0 \forall i \notin S$

Where  $r$  is the number of appearances of the related subjects in the page categories,  $ct$  is number of subjects that are a sub-string of the page name,  $k$  is a factor that can be customized in order to reach convergence (we chose 20 in our tests). and  $t = \sum_{i=0}^n P_0(i)$

Using this second method we get the following top and bottom results :

max 1 : Chess  
 max 2 : Grandmaster (chess)  
 max 3 : Boxing  
 max 4 : United States  
 max 5 : Federation Internationale des       checs  
 max 6 : Chess variant  
 max 7 : Chess Olympiad  
 max 8 : Chess problem  
 max 9 : Chess boxing  
 max 10 : World Chess Championship

min 1 : Pioneer High School  
 min 2 : Umea IK  
 min 3 : Guadalupe River (California)  
 min 4 : Brian Ashton (rugby player)  
 min 5 : Grammy Award for Best Latin Rock/Alternative Album  
 min 6 : Brampton Thunder  
 min 7 : Pump Up the Volume (song)  
 min 8 : Sigurd Evensmo  
 min 9 : Ethel Snowden  
 min 10 : St Edward's School (Oxford)

## 4 TME 3

**Exercise 1 - Simple Benchmark** Following are some of the generated random graphs :




FIGURE 7 – Random generated graph with  $p = 1.0$  and  $q = 0.0$




FIGURE 8 – Random generated graph with  $p = 0.4$  and  $q = 0.02$






FIGURE 9 – Random generated graph with  $p = 0.6$  and  $q = 0.1$




FIGURE 10 – Random generated graph with  $p = 0.7$  and  $q = 0.2$




FIGURE 11 – Random generated graph with  $p = 0.5$  and  $q = 0.5$

As we can see on the graphical rendition of the graphs, increasing  $\frac{p}{q}$  ratio leads to more noticeable separation in communities, while decreasing it creates a more homogeneous graph with only one big community.

**Exercise 2 - Label propagation** We have implemented the Label Propagation label in the `LabelPropagator` class. Following are the results of the call of this algorithm of some random generated graphs.



FIGURE 12 – Results of the label propagation algorithm on a random generated graph with  $p = 1.0$  and  $q = 0.0$




FIGURE 13 – Results of the label propagation algorithm on a random generated graph with  $p = 0.6$  and  $q = 0.01$




FIGURE 14 – Results of the label propagation algorithm on a random generated graph with  $p = 0.4$  and  $q = 0.02$




FIGURE 15 – Results of the label propagation algorithm on a random generated graph with  $p = 0.6$  and  $q = 0.1$




FIGURE 16 – Results of the label propagation algorithm on a random generated graph with  $p = 0.7$  and  $q = 0.2$




FIGURE 17 – Results of the label propagation algorithm on a random generated graph with  $p = 0.5$  and  $q = 0.5$

While the algorithm is efficient at recognizing communities with a  $\frac{p}{q}$  ratio higher than 20, when it decreases it has more trouble and only recognizes one unique community comprising the whole graph.

**Exercise 3 - New algorithm** For this matter, after searching for different existing algorithms, and in order to experiment on a different approach from label propagation and Louvain algorithms, we decided to turn to divisive approaches, and, after looking more in depth we settled on the Girvan-Newman algorithm. This algorithm computes the communities by successively removing edges from the graph by a certain measure, and then computing the connected components of the graph, which are the communities. The measure, called betweenness is defined as the extent to which that edge lies along shortest paths between all pairs of nodes. The edge with the highest betweenness will be removed at each iteration. We experimented on this algorithm using the `Jung` java API, in the `Clustering` class.